
Assignment A

Prefix / Suffix minima

Wang Yanhao - 3 July 2016

Introduction

Problem requirements:

Given an array $A = (a_1, \dots, a_N)$ with elements drawn from a linear ordered set. The suffix minima problem is to determine $\min\{a_i, a_{i+1}, \dots, a_N\}$ for each i . The prefix minima are $\min\{a_1, a_2, \dots, a_i\}$ for each i .

Assignment A requires us to design and implement an efficient parallel program with time complexity $O(\log(N))$ in both PThreads and OpenMP that computes the prefix and suffix minima of any given array. Moreover, an efficient sequential algorithm should be provided as well for later comparison of time taken between sequential implementation and parallel implementations, as described in Section Test Result.

Algorithm:

The intuitive sequential algorithm for prefix minima will have one for loop with i as the counter and check from the first element until the last element in the array. Another variable is needed to keep track of the current minima and update the prefix minima for each position. The sequential algorithm for suffix minima can be done similarly except this time the for loop counter goes from the index of the last element to the index of the first element.

However, since the value of the current minima depends on the position of the for loop counter, this sequential algorithm can not be easily paralleled. Hence, we need to use Balanced Tree method to implement a recursive algorithm where:

1. we generate an array consists of the minima of adjacent elements
2. we use a recursive call to compute the minima for the obtained array
3. we update the actual minima array from result generated above

Unlike the intuitive sequential algorithm, the first and third step for this Balanced Tree method can be parallelised easily.

Algorithm Proof

We only need to show the correctness for prefix minima because the proof of correctness for suffix minima is similar.

We prove the correctness of the algorithm by mathematical induction on k , where the size of the input is $n = 2^k$.

For prefix minima, when $k = 0$, the case is trivial because the result will be the only element itself.

Now assume the algorithm is true for some $k > 0$, we look at $n = 2^{k+1}$.

By induction, the result array after recursive step 2 holds the prefix minima of the sequence $(z_1, \dots, z_{N/2})$ where $z_i = \min(a_{2i-1}, a_{2i})$ for $1 \leq i \leq n/2$. Therefore, it is easy to see that when i is even, $\text{prefixmin}(a_i) = z_{i/2}$. When i is odd, assume $i = 2j + 1$, $\text{prefixmin}(a_i) = \text{prefixmin}(a_{2j+1}) = \min(\text{prefixmin}(a_{2j}), a_{2j+1}) = \min(z_{(i-1)/2}, a_i)$

Therefore, if the algorithm is true for some $k > 0$, it also works for the case of $k+1$.

Hence, we have proven the algorithm is true for all input size.

Analysis of time complexity

For step 1 and step 3, by using a parallelised implementation, the time complexity is $O(1)$ while the time complexity for step 2 is always halved. Therefore:

$$T(n) = T(n/2) + O(1)$$

and we can obtain the time complexity of the algorithm $T(n) = O(\log n)$.

Implementation

The sequential implementation is straightforward from the algorithm described in Section Introduction, however, there are a few points to take note of:

1. since prefix minima and suffix minima share the same part of algorithm (Step 1 and Step 2), we can shorten the code by using the same method with a flag to denote the if the current call is prefix minima or suffix minima.

2. Although the algorithm describe above assumes one-based array, the actual implementation use zero-based array, therefore, the odd / even index check in step 3 is reversed.

The OpenMP implementation comes directly from the sequential implementation by parallelising the for loop in Step 1 and Step 3.

The pthread implementation use similar strategy as the OpenMP implementation but requires more details. Before running the for loop, we decide the portion of indexes each available thread needs to work on by averaging (number of work / number of threads). And each thread will perform calculation in their assigned index range, which is not overlapping with other threads.

Test Result

Note:

The graphs generated in this section are not to scale due to the lack of time to write a proper code to generate graph, all credits of graphs goes to [onlinecharttool.com](https://www.onlinecharttool.com)

Environment:

The testing environment is the remote accessed computer via TUD277869.ws.tudelft.net

Result:

The raw test result generated by the code is shown below:


OpenMP:

NSize	Iterations	Seq	Th01	Th02	Th04	Th08	Par16
4096	1000	0.178122	0.183707	0.152617	0.113308	0.106043	0.134342
8192	1000	0.340917	0.371829	0.250169	0.182189	0.152370	0.181050
16384	1000	0.685468	0.742024	0.475054	0.316179	0.240619	0.255229
32768	1000	1.370638	1.487350	0.909530	0.591046	0.404177	0.400644
65536	1000	2.741190	2.959822	1.800727	1.115204	0.744628	0.647847
131072	1000	5.488728	5.915553	3.503353	2.148322	1.354764	1.156245
262144	1000	10.955003	11.802074	6.893868	4.165844	2.567574	2.163548

Pthread:

NSize	Iterations	Seq	Th01	Th02	Th04	Th08	Par16
4096	1000	0.169686	0.834247	0.943398	1.523053	4.266307	9.191767
8192	1000	0.357617	1.171309	1.172175	1.769501	4.744994	9.997933
16384	1000	0.690948	1.756396	1.474755	2.137263	5.301653	10.946510
32768	1000	1.371886	2.881639	1.975278	2.619067	5.943570	11.994456
65536	1000	2.745842	4.651124	2.989230	3.017233	6.924885	13.257850
131072	1000	5.492375	8.260639	5.414791	3.862719	8.462558	15.002484
262144	1000	10.960999	15.244151	9.840152	5.882813	10.789132	18.176377

The graph for T(N,P fixed at 2)

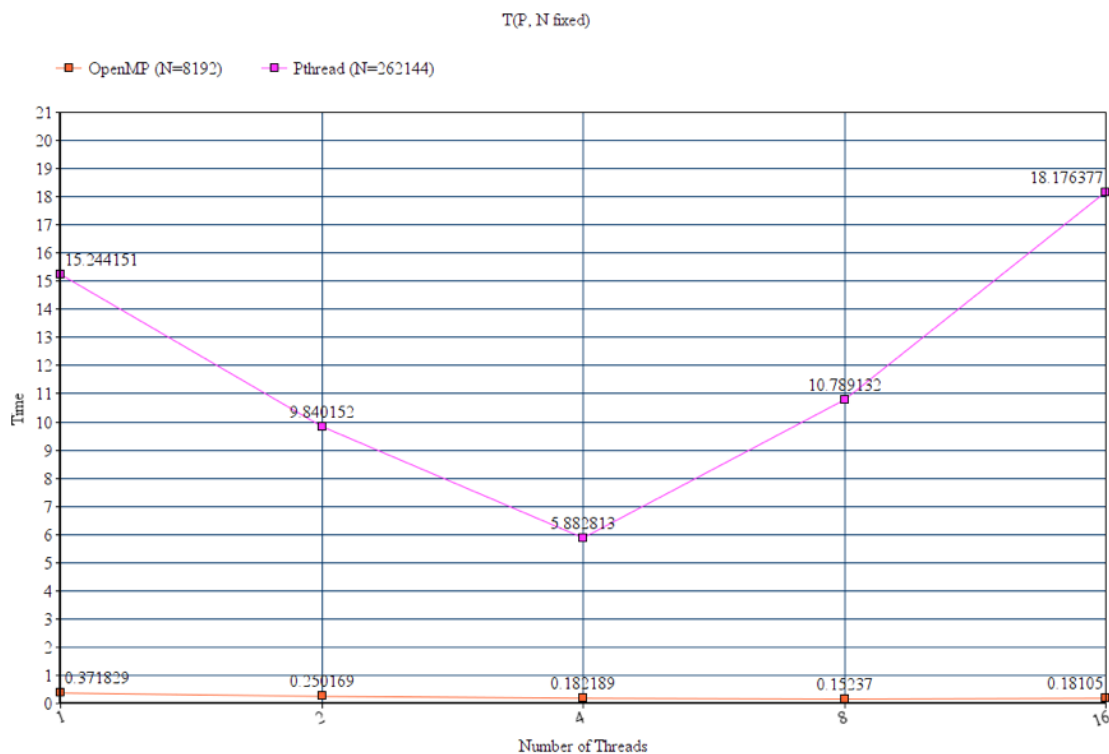


From the test result, it is obvious and intuitive to see that the time taken to execute the program increases with NSize.

Although the x-axis is not to scale, we can infer from the result that the time taken increases linearly as size of the input grow.

In addition, we can see that pthread implementation is slower than OpenMP implementation which might be caused by non-optimised pthread implementation.

Another interesting graph to see is the $T(P, N \text{ fixed})$:



Although the x-axis is not to scale, we can still see that there is a decrease when the number of threads increases, which is in accordance to our expectation.


However, the time taken to execute the program increases after certain point. This would be caused by the time overhead of creating and destroying threads. While concurrent execution can save us time, for small calculations, the communication overhead will outweigh the benefit gain if the number of threads go beyond the optimum number.

From the above analysis, we can see that for fixed input size, the time taken to execute the program will first decrease as the number of threads grow, then increase after the number of threads surpass a certain point.

Hence, it is intuitive to infer that, for fixed number of threads used in parallel execution, the ratio of the time taken for sequential execution over that for parallel execution will be small if the input size is small, and increase as the input size grow because there are more

work to work on, and the overhead of creating and destroying threads will in turn have less weightage.


The Speedup(N, P fixed at 2) graph shown below is in accordance with our intuition.



From the T(P, N fixed) graph we already observed that when input size is fixed at 262144, the time taken for the pthread implementation to execute the program will first decrease as the number of threads grow to 4, then increase afterwards. Similar curve reappears for OpenMP implementation when input size is fixed at 8192.

Since the Speedup is defined as the ratio of time taken for sequential execution over time taken for parallel execution, it is intuitive for us to deduce that for pthread implementation, when input size is fixed at 262144, the speedup will increase as the number of threads grow to 4, then decrease afterwards. And for OpenMP implementation, when input size is fixed at 8192, the speedup will increase as the number of threads grow to 8, then decrease afterwards.

Indeed, the Speedup(P, N fixed) graph shown in the next page is in accordance with our deduction. (only the Pthread data is plotted for clarity, but you can obtain similar curve when using data with N=4096 for OpenMP implementation)



Conclusion

In this exercise we have learnt to solving prefix minima and suffix minima problem using the Balanced Tree method and two ways to parallelise the sequential algorithm.

We have studied how to create multi-threaded programs in C programming language using OpenMP and pthreads libraries. It is much easier to use OpenMP to parallelise C programs whereas pthread grants us more control over the how each thread should perform.

From the test result, we can see that sequential algorithms generally perform better if the input size is small and not many calculation is involved whereas the benefit gained from concurrent execution will be surpassed by the overhead of creating and destroying threads when the input size is small.

However, if there are a lot of calculation needed for each thread to perform, parallelising the execution will grant us great performance benefits.

In practice, it might be useful to us to perform some experiment in order to estimate the communication overhead and determine the optimal number of threads to use for parallelising tasks.